

Tracking the Diffusion of Named Entities

TBD

Abstract

To do

Index Terms

To do

I. INTRODUCTION

The aim of this paper is to understand how named entities *emerge* and *spread* through social media based discourse. We are interested in exploring the following research questions:

- 1) **RQ1:** How can we accurately detect named entities in social media based discourse, given its myriad formats, often informal vernacular, and inherent noise (e.g. misspellings, abbreviations, etc.)?
- 2) **RQ2:** Under what conditions do entity mentions diffuse through discourse? And when are people *most likely* to be influenced into then discussing entities?
- 3) **RQ3:** How can we predict the discussion of certain named entities and who will begin talking about them?

II. DATASETS

For this research we will use the following datasets:

- 1) Reddit data – download and access all of the data from the full dump.¹
- 2) CoNLL 2003 data – a corpus of newswire texts, annotated for named entity chunks and types. This describes where entity mentions are in the text, including locations, organisations, and person mentions.
- 3) Twitter data; unannotated – we have a large corpus of English tweets that we can use here.
- 4) Twitter data; annotated – there are two datasets annotated with named entities. These are from Ritter’s 2011 EMNLP paper, and the W-NUT 2015 shared task.

III. RESEARCH STAGES

A. Stage 0: Data Preparation and NER

-To do:

- Annotate corpora with detected entities using basic typing of: person, location, organisation
- Run NER software over dataset and validate accuracy of this (using basic measures)
- Run NER over entire dataset to extract entities

¹https://archive.org/details/2015_reddit_comments_corpus

B. Stage 1: Exploratory Analysis

-To do:

- Plot relative frequency distribution as a function of time for named entities, and characterise the *shape* of the entities
- Apply lifecycle model to profile users' NER citations over time and investigate how users' profiles are influenced by global, community, and prior behaviour dynamics

C. Stage 2: Diffusion Analysis

-To do:

- Model the spread of named entities through user profiles (could use multivariate diffusion models here)

D. Stage 3: Forecasting

-To do:

- Implement models to forecast if a user will mention an entity and who that will be (hard!)

IV. RELATED WORK

twitter info propagation

reddit compared to other OSNs: "Lifespan and propagation of information in On-line Social Networks: A case study based on Reddit" <http://www.sciencedirect.com/science/article/pii/S1084804515001307>

network structure of reddit: "Navigating the massive world of reddit: using backbone networks to map user interests in social media" <https://peerj.com/articles/cs-4/>

twitter nlp

V. DATA PREPARATION AND NER

To conduct our study, we need to convert 140GB of compressed Reddit posts into a set of interlinked and time-ordered conversations and the entities mentioned in each of them. This provides a number of sub-challenges: sampling of the Reddit data, creating a linked series of conversations, and picking out entity mentions in this text type. Reddit data is largely unexplored in the NLP community, despite the large volume of it and the especially rich metadata. This poses additional challenges: certainly, given the lack of work on Reddit text, there are no annotated datasets available yet, so supervised in-domain work is not directly possible yet. Additionally, the datasets are large, which makes it important to choose a good subset of data on which to do prototyping and development, in order to keep research cycles short. The result that we come to at the end of this stage is a rich dataset for tracking entity and concept diffusion within and across communities.

The Reddit dataset [1] is comprised of a sequence of comments, with one JSON record for each one. These are ordered temporally. Reddit itself is roughly similar to a forum, where top-level divisions are made by topic. Within each topic, or *subreddit*, there are posts, which begin with either a short piece of text or a link to an external resource – typically an image, video, or interesting article. Users then may publish comments for each post, and



Fig. 1. Example discussion around a Reddit post.

reply to each others' comments. This leads to a threaded discussion, centred on a particular topic, with a hierarchical comment structure (see Figure V).

A. Subreddit extraction

- subreddit extraction
- top lists
- describe top lists site
- pick 100
- extract these over whole sample
- figures for raw dataset (# per year; total volume over time; volume-rank graph of subreddits; possible ref to appx)

B. NER for Reddit

- where is NER first mentioned and defined?

We model micro-topics in conversation as entity mentions. This allows tracking of topics at maximally fine granularity, looking at each user's interests at a low level, as opposed to monitoring broader topics such as "consumer electronics", "politics" and so on. In fact, these broader topics are already explicitly annotated by means of the subreddit topics.

Entity mentions are extracted through named entity recognition. Generally, this task aims to detect the boundaries of certain kinds of entities within a certain piece of text. In this instance, we tokenise text, splitting it into

sentences using the Punkt tokeniser [2], and subsequently word-sized chunks, using the `twokenize` tool with some adaptations [3]. This tool performs Penn Treebank-style tokenisation, a common standard, with some specific adaptations to enable it to handle the noise present in user-generated text. After this, we take a structured prediction approach to deciding which tokens in each sentence are part of an entity, and possibly the type of the entity. Finally, we concatenate entity tokens, and use these to build a list of entity mentions in any given input text. For example, given the input comment from the source JSON:

“body”: “There are still some really good fighters on this card. Conor McGregor is on the card and so is Gunnar Nelson.”

The following output entities should be collected:

“entity_texts”: [“Conor McGregor”, “Gunnar Nelson”]

Typically, many NER systems take a supervised approach; that is, they use data labelled by humans as training data, from which features are extracted to form training instances for a machine learning algorithm. However, NLP systems can be hard to transfer between text types; for example, NER systems for newswire might reach about 89% F1 on news articles, but only around 40% on tweets (a form of user-generated content), as found in [4]. One approach to overcoming this performance drop when changing text type is to train over a blend of text types. For example, Ritter [5] used both IRC² and newswire data when developing a part-of-speech tagger for tweets, as well as an unsupervised language model from the target text type. This led to strong performance improvements. We follow a similar approach, using a blend of NE-annotated corpora from both newswire and tweets. The newswire data is drawn from the CoNLL-2003 evaluation task set [6]; the twitter data is from Ritter’s early experiments and also the W-NUT 2015 shared task [5], [7].

We start using structured predicting in the form of a CRF to label whole sentences at a time. For features, we use a fairly classical set, and add some unsupervised word representations to this. Our base features are:

-

In addition, we induce Brown clusters [8] and use these as word representations [9]. Brown clustering is a form of hierarchical agglomerative hard clustering, using average mutual information as its objective function. It takes as input a corpus, in the form of a sequence of words, and in its generalised form [10], a single hyperparameter: the size of its active set a . The result is a sequence of binary merges, describing the set membership of each word type in the corpus as the merges progress. For each single word type, therefore, the path to a destination cluster can be described as a bitstring, which details the sequence of binary merges taken. The zero-length bitstring describes the situation at the top of the hierarchy, where there is one class.

These bitstrings are typically converted to features by shearing. This involves only examining the first n bits of a bitstring. However, shearing does not maximise the information preserved in the representation – sub-clusterings at many levels are lost. We therefore experiment with a new method of feature extraction from Brown clusters, which we call *provenance-based* feature extraction. We take the cluster identifier at every level, tracing the provenance of a terminal word cluster all the way to the root cluster (which contains all word types). This preserves the entire set

²Internet Relay Chat – informal internet conversation text

membership of any given term, throughout the induced hierarchical clustering. As an example, if the term *fishing* has bitstring 1100101, the following text features are generated: 1, 11, 110, 1100, 11001, 110010, 1100101. For contrast, if the typical bit depths of 4, 6, 10 and 20 [11] were chosen, only the following features would be generated: 1100, 110010, 1100101. As a result of taking all directly-relevant features in the merge list, the lossy nature of shearing-based feature extraction from Brown clusters is avoided.

Feature extraction, training, classification and JSON annotation are all performed using an entity recognition toolkit (https://github.com/leondz/entity_recognition [12]), with custom extractors.

C. Tuning entity recognition

Entity recognition needs to be tuned to fit Reddit data well. There are a number of parameters in our training data balance, feature extraction, and objective function that all reflect the nature of the data and the task at hand. We present our method for estimating of these factors, and intrinsic NER evaluation.

In terms of evaluation, we prefer recall over precision. Over the large dataset, spurious entities are likely to be those that are not seen very often or have an unusual pattern. This suggests that there will be great variation in their surface forms, leaving them in the long tail of discovered entities. As the majority of our work looks at the more frequent patterns, these are less likely to have an impact. Conversely, recall reveals how well the extraction is performing, and it important to track a range of entities. In addition, recall has always been more challenging to achieve in social media texts than recall [5], [4].

To capture this balance, we take F-scores with $\beta = 2$. Given precision P and recall R , typically an F-score is drawn from F_β with $\beta = 1$.

$$F_\beta = (1 + \beta^2) \frac{PR}{(\beta^2 P) + R} \quad (1)$$

When $\beta = 1$, precision and recall are balanced in a harmonic mean, e.g. F1-score. That is, false positives and false negatives impact results equally. To score away from false positives, i.e. wrong cues, we set $\beta = 2$.

Our approach here is to tune an entity recogniser with reference to a dataset that matches the target text type. We draw this development set from Reddit posts, using comments encountered during our work that appear to have missing or spurious annotations. These are then isolated, tokenised, and manually annotated. Our annotation format follows [5] in using the Freebase top-level entity type inventory, but only uses the chunking information, as nothing further than this is needed: only the surface forms of entities. In total, we identified and annotated 3 708 tokens of Reddit data, including 149 entity chunks. This comprised our development set, which was used to tune a variety of parameters in our approach.

Firstly, we tuned our word representations. Specifically, we needed to estimate the number of Brown clusters C to use in feature extraction. We also then used this to determine what blend of social, Reddit or news data gave best results in unsupervised feature generation and extraction. To tune the value for C , we examined a similar scenario with similar dataset sizes, and estimated an optimum. We noted that in prior work [13], entity recognition performance with decomposed class prefixes – similar to the provenance-based features we propose here – peak at

around $C = 2500$ for corpora of 16k tokens, $C = 5000$ for corpora of 32k tokens, and at higher values for larger datasets. As General Brown clustering is dependent on the number of types and the size of the active set a , and results are unreliable with $a > C$, we set $C = a = 4000$. We then experiment with combinations of newswire, Twitter and Reddit data. Brown clusters are extracted using the generalised-brown package [14]. Results are given in Table ??.

– results of brown cluster tuning

In addition, we draw supervised data for multiple datasets in order to approximate the Reddit text type. We take data from Reddit, taking the union of corpora used in previous work that follow the Freebase ten-class entity scheme [5], [7]. This scheme gives broader coverage than e.g. the three-class ACE named entity scheme. For newswire, we use the Reuters RCV1 corpus annotations that were part of the CoNLL-2003 shared task [6]. Classes are removed before training, making this just a chunking task. We compare against Stanford NER [15] as a baseline. Results for different balances of training data are given in Table ??

– analysis of cross-genre ne chunking performance

Note that while some figures seem low when compared to typical newswire level performance, the toolkit used is high-ranking, state-of-the-art research software, coming third in the 2015 W-NUT challenge for entity chunking over tweets. The task is simply difficult; Twitter NER recall has always been low. In addition, it is a generally consistent finding that generalising NER systems beyond newswire is not yet well understood; systems that perform very well on this text type (e.g. F1 of 0.89 from Stanford NER [4]) can often score very poorly on social media content (F1 of 0.41). This may be due to overfitting of tools to newswire over time, due to community challenges, dataset in just one type, extra custom rules adapting to formal news text, or other things – but this is beyond the scope of this paper. We do note that our approach uses largely unsupervised feature extraction and performs better than on other social media corpora, also beating the Stanford NER system, in this first attempt at named entity chunking for Reddit.

VI. ENTITY DIFFUSION

A. Entity Mention Cascades

-Describe the entity cascade graphs and what these entail -Explain how these were derived using graph isomorphism and then the rank induced

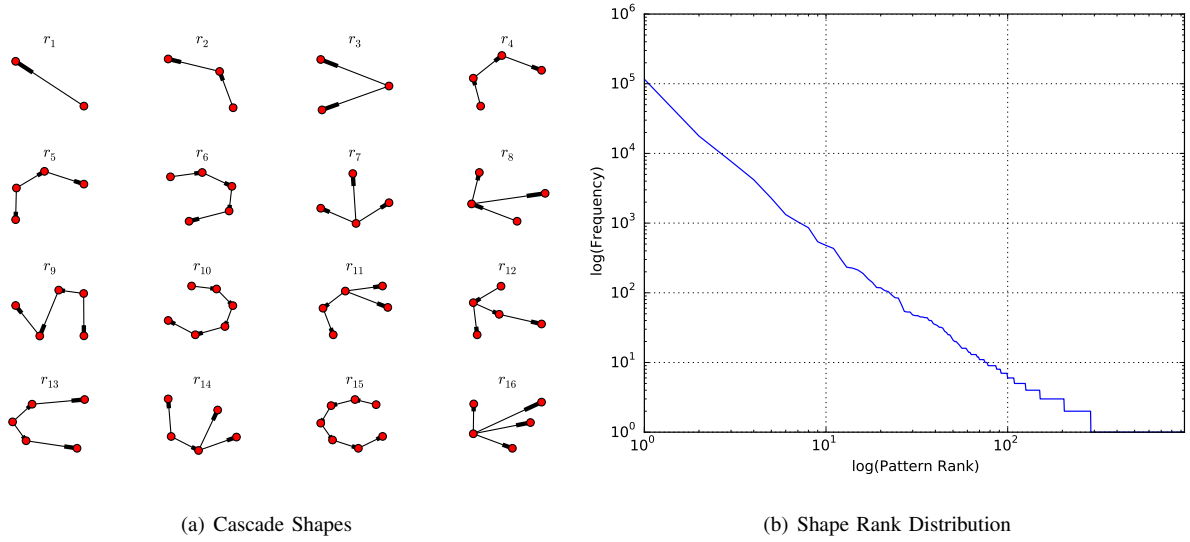


Fig. 2. To do

B. Entity Adoption Post- $k - 1^{th}$ Exposure

-Potential for adoption at the point of one exposure -Per entity-exposure dynamics - add plots of random selection of entities

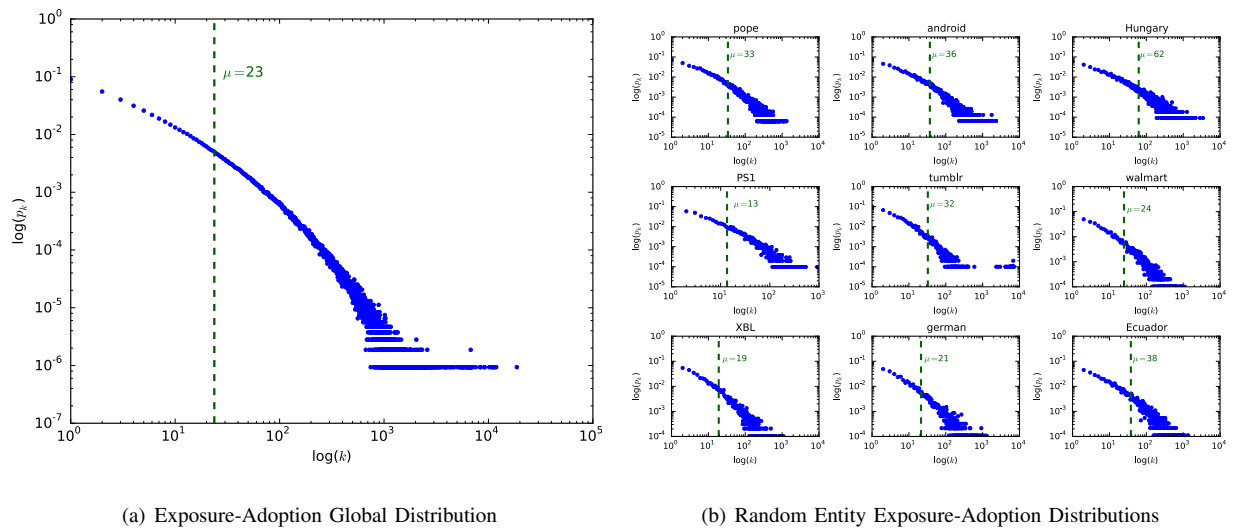


Fig. 3. To do

C. Global Threshold Diffusion Model

-Explain the implementation of WSDM2010 work and adaptation for non-sequential dataset scans - necessary due to memory load being infeasible

1) *Distribution of $\tau_{v,u}$* : -Present the distribution of $\tau_{v,u}$ and what this implies for the model




Fig. 4. To do

D. Experiments

-Explain the evaluation process and the computation of the ROC values - per entity -Plot variance in the ROC values achieved

REFERENCES

- [1] J. Baumgartner, "Complete public Reddit comments corpus," 2015.
- [2] T. Kiss and J. Strunk, "Unsupervised multilingual sentence boundary detection," *Computational Linguistics*, vol. 32, no. 4, pp. 485–525, 2006.
- [3] B. O'Connor, M. Krieger, and D. Ahn, "TweetMotif: Exploratory Search and Topic Summarization for Twitter," in *Proc. ICWSM*. AAAI, 2010.
- [4] L. Derczynski, D. Maynard, G. Rizzo, M. van Erp, G. Gorrell, R. Troncy, J. Petrak, and K. Bontcheva, "Analysis of named entity recognition and linking for tweets," *Information Processing & Management*, vol. 51, no. 2, pp. 32–49, 2015.
- [5] A. Ritter, S. Clark, O. Etzioni *et al.*, "Named entity recognition in tweets: an experimental study," in *Proceedings of the Conference on Empirical Methods in Natural Language Processing*. ACL, 2011, pp. 1524–1534.
- [6] E. F. Tjong Kim Sang and F. De Meulder, "Introduction to the CoNLL-2003 shared task: Language-independent named entity recognition," in *Proceedings of the seventh Conference on Natural Language Learning*. ACL, 2003, pp. 142–147.
- [7] T. Baldwin, Y.-B. Kim, M. C. de Marneffe, A. Ritter, B. Han, and W. Xu, "Shared tasks of the 2015 workshop on noisy user-generated text: Twitter lexical normalization and named entity recognition," *Proc. W-NUT (ACL-IJCNLP)*, pp. 126–135, 2015.
- [8] P. F. Brown, P. V. Desouza, R. L. Mercer, V. J. D. Pietra, and J. C. Lai, "Class-based n-gram models of natural language," *Computational linguistics*, vol. 18, no. 4, pp. 467–479, 1992.
- [9] J. Turian, L. Ratinov, Y. Bengio, and D. Roth, "A preliminary evaluation of word representations for named-entity recognition," in *NIPS Workshop on Grammar Induction, Representation of Language and Language Learning*, 2009, pp. 1–8.

- [10] L. Derczynski and S. Chester, “Generalised Brown Clustering and Roll-Up Feature Generation,” in *Proceedings of the 30th conference of the Association for Advancement of Artificial Intelligence*, 2016.
- [11] L. Ratinov and D. Roth, “Design challenges and misconceptions in named entity recognition,” in *Proceedings of the Thirteenth Conference on Computational Natural Language Learning*. Association for Computational Linguistics, 2009, pp. 147–155.
- [12] L. Derczynski, I. Augenstein, and K. Bontcheva, “USFD: Twitter NER with Drift Compensation and Linked Data,” *Proc. W-NUT (ACL-IJCNLP)*, pp. 48–53, 2015.
- [13] L. Derczynski, S. Chester, and K. S. Bøgh, “Tune Your Brown Clustering, Please,” in *Proceedings of the conference on Recent Advances in Natural Language Processing (RANLP)*, 2015.
- [14] S. Chester and L. Derczynski, “generalised-brown: Source code for AAAI 2016 paper,” Nov. 2015. [Online]. Available: <http://dx.doi.org/10.5281/zenodo.33758>
- [15] J. R. Finkel, T. Grenager, and C. Manning, “Incorporating non-local information into information extraction systems by gibbs sampling,” in *Proceedings of the 43rd Annual Meeting on Association for Computational Linguistics*. Association for Computational Linguistics, 2005, pp. 363–370.
- [16] L. Derczynski, D. Maynard, N. Aswani, and K. Bontcheva, “Microblog-genre noise and impact on semantic annotation accuracy,” in *Proceedings of the 24th ACM Conference on Hypertext and Social Media*. ACM, 2013, pp. 21–30.